

Alexey Yurchenko

AI Agent Engineer / Senior Fullstack Engineer

Email: alexes.dev@gmail.com

GitHub: github.com/alexesdev

Remote / Digital nomad

ABOUT ME

I'm looking for an **AI Agent Developer** role. Most of the code I write these days isn't for humans. It's for agents, and the agents are the ones working with humans. The long-term goal is to run evolutionary selection inside swarms of agents. I'm not there yet — I'm building the tools first.

15 years of practice as a developer, since 2011, without a week away from code. Right now I run iiminion.ru, where I package autonomous AI agents and sell them by subscription: an agent that watches competitors and writes posts in a brand voice, a script-to-video pipeline, a landing-page generator that survives client edits, plus per-customer fine-tuning. The infrastructure underneath it is mine — an agent-orchestration platform (Hermes on a Nomad cluster, agents coordinating through shared memory) and an open-source MCP knowledge protocol called trip2g.com.

Python and Go are my main languages. Most of my career has been about taking unclear specs and turning them into systems that actually run. I've gotten good at working without a plan when there isn't one.

FEATURED PROJECTS

AI Agent Orchestration Platform — Independent R&D

2025 — Present

Infrastructure for autonomous AI agents that operate on a shared, transparent knowledge substrate.

- Built a control panel for deploying Hermes agents into a Nomad cluster — one-click rollout, lifecycle management, fleet-level observability.
- Packaged the agent so it maintains its own Obsidian vault as a "second brain" — every action is recorded with its rationale (what was done, how, why), giving the agent transparent, human-readable memory and full auditability.
- Multi-agent coordination through shared memory: agents read each other's vaults to align on context, hand off tasks, and avoid duplicate work — a lightweight alternative to heavy orchestration frameworks.
- Designed to interoperate with trip2g.com so agent vaults are indexable and queryable across the same federated MCP protocol.

trip2g.com — Founder, Sole Engineer (open source, MIT)

2025 — Present

A federated protocol that turns markdown note collections (Obsidian, Notion, Google Drive, Telegram) into RAG sources for AI agents over Model Context Protocol. About 12 months of solo work from the first prototype to a public hub where other people now plug in their own bases.

It didn't start as an AI project. I'd been keeping an Obsidian vault for three years and wanted a way to share parts of it. Writing long posts is hard; writing 5–10 short interlinked notes is much easier. The original plan was a knowledge-graph engine where you could carve out a subgraph and sell it. I knew AI should fit in somewhere but couldn't see how. Once I tried MCP, the picture clicked: the base

becomes a private external RAG, the sensitive data stays on the user's machine, and the agent — not the human — is the natural consumer. Everything below (federation, the vector router, the self-describing bases) is what I had to build to make that workable. The same realization is what spun off iiminion.ru.

- **Markdown-first knowledge engine** — ingests markdown, renders each note as a page, exposes the corpus as a typed knowledge graph; arbitrary schemas and monetizable subgraphs can be built on top.
- **MCP server as external RAG** — keeps sensitive data local; the agent issues targeted queries instead of shipping the full corpus to a third-party service.
- **Self-describing bases** — each base ships its own search instructions, letting the calling agent *plan its retrieval* rather than rely on a fixed retriever.
- **Vector "router" base** — a meta-base over other bases: semantic search routes a query to the right downstream sources, then the agent fans out across the matches.
- **Protocol-level federation with adapters** for Notion, Google Drive, Telegram, etc. — direct hub-to-hub peer-to-peer, no central relay.
- **Recent features:** per-heading note retrieval and TOC over MCP; admin-mode API key (the agent can manage the system itself); BEM templating that makes agent-driven landing-page edits structurally stable.
- External users are building their own subgraphs and case-study bases on the protocol.

WORK EXPERIENCE

Founder & AI Engineer — iiminion.ru

2026 — Present | Remote

Productizing autonomous AI agents on a subscription basis; agents are packaged for self-serve use and fine-tuned per customer to match their workflow.

- **Competitor-monitoring agent:** tracks competitors and produces branded social posts and content for the customer's product in their voice — research, drafting, and publishing.
- **Script-to-video pipeline:** end-to-end automated short-form video generation from a written script.
- **Edit-stable landing-page generator:** produces marketing landing pages that survive iterative client edits without structural breakage.
- Subscription delivery model; per-customer fine-tuning for niche workflows.
- Built on top of the agent-orchestration platform and MCP protocol described above — own infrastructure end-to-end.

Backend Developer — P2P.org (P2P Validator)

08/2022 — 01/2024 | Remote

Infrastructure and data platform for Proof-of-Stake blockchain networks.

- Built data extractors in **Python and Go** for Cosmos SDK-based networks; pipelines processed terabytes of on-chain data into BigQuery for the analytics team.
- Shipped APIs over gRPC + Protobuf and React / TypeScript dashboards used daily by internal data analysts.
- Owned end-to-end pipelines: Kafka ingestion → transformations → BigQuery → analyst-facing dashboards.
- Deployed on Google Cloud / Kubernetes via Docker; collaborated cross-team on schema design and data SLAs.

Stack and tools: Python, Go, React, TypeScript, BigQuery, Kafka, gRPC / Protobuf, Docker, Kubernetes, Google Cloud, Cosmos SDK, Cosmos Hub, Agoric.

Head of IT Department — FRESH / MoscowFresh

07/2018 — 04/2022 | Moscow

Brought back to FRESH after a short stint at Statsbot specifically to lead the technical integration of the company's first acquisition. Stayed on to run the entire in-house technology function during a high-growth, M&A-heavy phase — 2 Moscow warehouses, ~1,000 orders/day.

- Managed a cross-functional team of 10 (backend, frontend, mobile, design, PM); led hiring, mentoring, and technical direction.
- **Led technical integration of acquired competitors** — consolidated their stacks, data, and operations into FRESH systems across multiple M&A events.
- **Architected and built the warehouse management system from scratch** — stock, supplies, shipments, full warehouse-accounting automation — replacing manual operations across both sites.
- **Implemented company-wide business analytics on Metabase**, building on data-analytics experience from Statsbot; gave every department self-serve dashboards on top of operational data.
- Shipped role-specific tools for every department: drivers, logistics managers, cooks, packers, assemblers; integrated the WMS with industrial cash registers and weighing scales over the network (binary protocol).
- Built the customer-facing storefronts (Gatsby / React) and React Native mobile apps; backend on Go + [gqlgen](#) (GraphQL).
- Operated a 30+ VM cluster on Yandex Cloud with Nomad / Consul, Packer / Ansible provisioning, Prometheus / Grafana / AlertManager monitoring, and self-hosted GitLab CI/CD.
- Owned the full PoC → production cycle for every system shipped, often translating vague business requests directly into working tooling.

Stack and tools: Go, gqlgen, GraphQL, React, Gatsby, React Native, TypeScript, PostgreSQL, PostgREST, Postgraphile, Metabase, Nomad, Yandex Cloud, Ansible, Packer, Docker, Prometheus, Grafana, GitLab, Redis.

Fullstack Developer — Statsbot.co

12/2017 — 04/2018 | Moscow

Brief, focused stint at a SaaS data-analytics startup (the codebase later became the open-source project **Cube.js**) between FRESH tenures.

- Shipped features on Ruby on Rails / React in a fast-moving early-stage team.
- **First deep exposure to business analytics and OLAP workflows** — directly led to designing and implementing the company-wide analytics platform after returning to FRESH.

Stack and tools: Ruby, Ruby on Rails, React, Redux, Webpack, Docker.

Fullstack Developer — FRESH / MoscowFresh

01/2014 — 12/2017 | Moscow

Same company as the later Head of IT role — joined as the sole engineer and grew with the business.

- Took over and stabilized the e-commerce platform (Rails / React) after the previous developer; built the mobile web with SSR.
- Maintained the native iOS app (Objective-C) and migrated it to a React Native rewrite.

Stack and tools: Ruby, Ruby on Rails, React, SSR, React Native, Redux, Webpack, PostgreSQL, Redis, Sidekiq.

Freelance Web Developer

2011 — 2014 | Sevastopol

- Built websites and web apps for SMB clients in PHP / Yii and Ruby on Rails; HTML / CSS / jQuery frontends; PostgreSQL and MongoDB backends.

Stack and tools: PHP, Yii, Ruby, Ruby on Rails, HTML, CSS, jQuery, PostgreSQL, MongoDB.

SKILLS

- **AI / LLM:** OpenAI API, Model Context Protocol (MCP), AI agents, vector indexing, RAG, prompt engineering, agent orchestration, fine-tuning
- **Languages:** Python*, Go (Golang)*, TypeScript, Ruby*
- **Backend & Data:** REST, GraphQL, gRPC / Protobuf, Kafka, BigQuery, PostgreSQL*, Redis
- **Cloud & Infra:** Docker*, Kubernetes, Google Cloud, S3, Nomad / Consul, Ansible, Packer
- **Frontend:** React*, Next.js, Gatsby, React Native*, Tailwind

* — Confident with 5+ years of commercial experience

LANGUAGES

- Russian — Native
- Ukrainian — Native
- English — B2+

HOBBIES

Zettelkasten / Second Brain

3+ years on a personal Obsidian graph. Directly seeded trip2g.com.

Electronics

Designing and assembling small devices on popular microcontrollers. Useful skills: C++.

3D Printing

Designing parts in CAD and printing them. Useful skills: CAD proficiency, CNC experience.

Hydroponics

Soil-free plant growth with automation. Useful skills: nutrition / environment monitoring on Prometheus / Grafana.

GameDev

Unity, Godot, VR. Useful skills: C#, C++.